Breaking the Nonsmooth Barrier: A Scalable Parallel Method for Composite Optimization

Fabian Pedregosa^{†‡}, Rémi Leblond[†], Simon Lacoste–Julien* [†]INRIA and École Normale Supérieure, Paris, France. [‡]Currently at UC Berkeley *MILA and DIRO, Université de Montréal, Canada











Summary

Optimization methods need to be adapted to the **parallel** setting to leverage modern computer architectures.

Highly efficient variants of stochastic gradient descent have been recently proposed, such as Hogwild [1], Kromagnon [2], ASAGA [3].

They assume that the objective function is smooth, so are inapplicable to problems such as Lasso, optimization with convex constraints, etc.

Main contributions:

- 1. Sparse Proximal SAGA, a sparse variant of the linearly-convergent proximal SAGA algorithm.
- 2. **ProxASAGA**, the first parallel asynchronous variance-reduced method that supports *nonsmooth* composite objective functions.

Problem Setting

Objective: develop parallel asynchronous method for problems of the form

$$\underset{\boldsymbol{x} \in \mathbb{R}^p}{\text{minimize}} \frac{1}{n} \sum_{i=1}^n f_i(\boldsymbol{x}) + h(\boldsymbol{x}),$$

- f_i is differentiable with L-Lipschitz gradient.
- h is block-separable $(h(x) = \sum_B h_B([x]_B))$ and "simple" in the sense that we have access to

$$\mathbf{prox}_{\gamma h} \stackrel{\mathsf{def}}{=} \mathbf{arg\,min}_{\boldsymbol{x}} \gamma h(\boldsymbol{x}) + \frac{1}{2} ||\boldsymbol{x} - \boldsymbol{z}||^2$$
.

includes Lasso, group Lasso or ERM with box constraints.

Variance-reduced stochastic gradient methods are natural candidates due to their state of the art performance and recent asynchronous variants.

The **SAGA** algorithm [4] maintains current iterate $\boldsymbol{x} \in \mathbb{R}^p$ and historical gradients $\boldsymbol{\alpha} \in \mathbb{R}^{n \times p}$. At each iteration, sample $i \in \{1, \dots, n\}$ and compute $(\boldsymbol{x}^+, \boldsymbol{\alpha}^+)$ as

$$\boldsymbol{x}^+ = \mathbf{prox}_{\gamma h}(\boldsymbol{x} - \gamma(\nabla f_i(\boldsymbol{x}) - \boldsymbol{\alpha}_i + \overline{\boldsymbol{\alpha}})); \ \boldsymbol{\alpha}_i^+ = \nabla f_i(\boldsymbol{x}).$$

Difficulty of a Composite Extension

- Existing methods exhibit best performance when updates are sparse.
- Even in the presence of sparse gradients, the SAGA update is not sparse due to the presence of $\overline{\alpha}$ and prox.
- Existing convergence proofs bound noise from asynchrony using the Lipschitz constant of the gradient. This property does not extend to composite case.

A New Sequential Algorithm: Sparse Proximal SAGA

The algorithm relies on the following quantities

- Extended support T_i : set of blocks that intersect with ∇f_i . $T_i \stackrel{\mathsf{def}}{=} \{B : \mathsf{supp}(\nabla f_i) \cap B \neq \varnothing, \ B \in \mathcal{B}\}$
- For each block $B \in \mathcal{B}$, $d_B \stackrel{\text{def}}{=} n/n_B$, where $n_B := \sum_i \mathbb{1}\{B \in T_i\}$ is the number of T_i that contain B.
- $m{D}_i$ is a diagonal matrix defined block-wise $[m{D}_i]_{B,B} \stackrel{\mathsf{def}}{=} d_B \mathbb{1}\{B \in T_i\} m{I}_{|B|}.$
- φ_i is a block-wise reweighting of h: $\varphi_i \stackrel{\mathsf{def}}{=} \sum_{B \in T_i} d_B h_B({m x})$

Justification. The following properties are verified

 $arphi_i(m{x})$ is zero outside T_i $D_im{x}$ is zero outside T_i (sparsity) $\mathbb{E}_i\,arphi_i=h$ $\mathbb{E}_i\,m{D}_i=m{I}$ (unbiasedness)

Algorithm. As SAGA, it maintains current iterate $\boldsymbol{x} \in \mathbb{R}^p$ and table of historical gradients $\boldsymbol{\alpha} \in \mathbb{R}^{n \times p}$. At each iteration, it samples an index $i \in \{1, \dots, n\}$ and computes next iterate $(\boldsymbol{x}^+, \boldsymbol{\alpha}^+)$ as

$$egin{aligned} oldsymbol{v}_i &=
abla f_i(oldsymbol{x}) - oldsymbol{lpha}_i + oldsymbol{D}_i \overline{oldsymbol{lpha}} \ oldsymbol{x}^+ &= \mathbf{prox}_{\gamma arphi_i} \left(oldsymbol{x} - \gamma oldsymbol{v}_i
ight) \; ; \; oldsymbol{lpha}_i^+ &=
abla f_i(oldsymbol{x}) \end{aligned}$$

Features

- Per Iteration cost in $\mathcal{O}(|T_i|)$.
- Easy to implement (compared to the lagged update approach [5]).
- Amenable to parallelization.

Convergence Analysis

For step size $\gamma = \frac{1}{5L}$ and f μ -strongly convex ($\mu > 0$), Sparse Proximal SAGA converges geometrically in expectation. At iteration t we have

$$\mathbb{E}\|\boldsymbol{x}_t - \boldsymbol{x}^*\|^2 \le (1 - \frac{1}{5}\min\{\frac{1}{n}, \frac{1}{\kappa}\})^t C_0$$

with $C_0 = \|\boldsymbol{x}_0 - \boldsymbol{x}^*\|^2 + \frac{1}{5L^2} \sum_{i=1}^n \|\boldsymbol{\alpha}_i^0 - \nabla f_i(\boldsymbol{x}^*)\|^2$ and $\kappa = \frac{L}{\mu}$ (condition number).

Implications

- Same convergence rate than SAGA with cheaper updates.
 - In the "big data regime" $(n \ge \kappa)$: rate in $\mathcal{O}(1/n)$.
 - In the "ill-conditioned regime" $(n \le \kappa)$: rate in $\mathcal{O}(1/\kappa)$.
- Adaptivity to strong convexity, i.e., no need to know strong convexity parameter to obtain linear convergence.

A New Parallel Algorithm: Proximal Asynchronous SAGA (ProxASAGA)

Proximal Asynchronous SAGA (ProxASAGA) runs Sparse Proximal SAGA asynchornously and without locks and updates x, α and $\overline{\alpha}$ in shared memory.

All read/write operations to shared memory are *inconsistent*, i.e., no vector-level locks while reading/writing.

1: keep doing in parallel

```
Sample i uniformly in \{1, ..., n\}
                  [\hat{\boldsymbol{x}}]_{T_i} = \text{inconsistent read of } \boldsymbol{x} \text{ on } T_i
                 \hat{oldsymbol{lpha}}_i= inconsistent read of oldsymbol{lpha}_i
                  [\overline{\alpha}]_{T_i} = \text{inconsistent read of } \overline{\alpha} \text{ on } T_i
                  [\delta oldsymbol{lpha}]_{S_i} = [
abla f_i(\hat{oldsymbol{x}})]_{S_i} - [\hat{oldsymbol{lpha}}_i]_{S_i}
                 [\,\hat{oldsymbol{v}}\,]_{T_i} = [\deltaoldsymbol{lpha}\,]_{T_i} + [\,oldsymbol{D}_i\overline{oldsymbol{lpha}}\,]_{T_i}
                  [\delta oldsymbol{x}]_{T_i} = [\mathbf{prox}_{\gamma_{\mathcal{O}_i}}(\hat{oldsymbol{x}} - \gamma \hat{oldsymbol{v}})]_{T_i} - [\hat{oldsymbol{x}}]_{T_i}
                 for B in T_i do
                           for b in B do
                                       [\boldsymbol{x}]_b \leftarrow [\boldsymbol{x}]_b + [\delta \boldsymbol{x}]_b
                                                                                                                               if b \in \text{supp}(\nabla f_i) then
                                                 [\overline{\boldsymbol{\alpha}}]_b \leftarrow [\overline{\boldsymbol{\alpha}}]_b + 1/n[\delta \boldsymbol{\alpha}]_b
                                                                                                                               > atomic
                                       end if
                           end for
                end for
                oldsymbol{lpha}_i \leftarrow 
abla f_i(\hat{oldsymbol{x}})
                                                        (scalar update)
                                                                                                                               > atomic
18: end parallel loop
```

Perturbed Iterate Framework

Problem: Analysis of asynchronous parallel algorithms is *hard*.

Solution: Cast them as sequential algorithms working on *perturbed* inputs. Distinguish:

- \hat{x}_t : inconsistent vector. Counter t is incremented when a core *finishes reading* the parameters (after read labeling [3]).
- $m{x}_t$: the *virtual* iterate defined by $m{x}_{t+1} \stackrel{\text{def}}{=} m{x}_t \gamma m{g}_t$ with $m{g}(m{x}, m{v}, i) = rac{1}{\gamma} (\hat{m{x}}_t \mathbf{prox}_{\gamma arphi_i} (\hat{m{x}}_t \gamma \hat{m{v}}_{i_t}))$.

Interpret \hat{x}_t as a noisy version of x_t due to asynchrony. Generalization of perturbed iterate framework [2, 3] to composite objectives.

Analysis preliminaries

Definition (measure of sparsity). Let $\Delta := \max_{B \in \mathcal{B}} |\{i : T_i \ni B\}|/n$. This is the normalized maximum number of times that a block appears in the extended support. We always have $1/n \le \Delta \le 1$.

Definition (delay bound). τ is a uniform bound on the maximum delay between two iterations processed concurrently.

Convergence guarantee of ProxASAGA

Suppose $\tau \leq \frac{1}{10\sqrt{\Delta}}$. Then:

- If $\kappa \geq n$, then with step size $\gamma = 1/36L$, ProxASAGA converges geometrically with rate factor $\Omega(\frac{1}{\kappa})$.
- If $\kappa < n$, then using the step size $\gamma = 1/36n\mu$, ProxASAGA converges geometrically with rate factor $\Omega(\frac{1}{n})$.

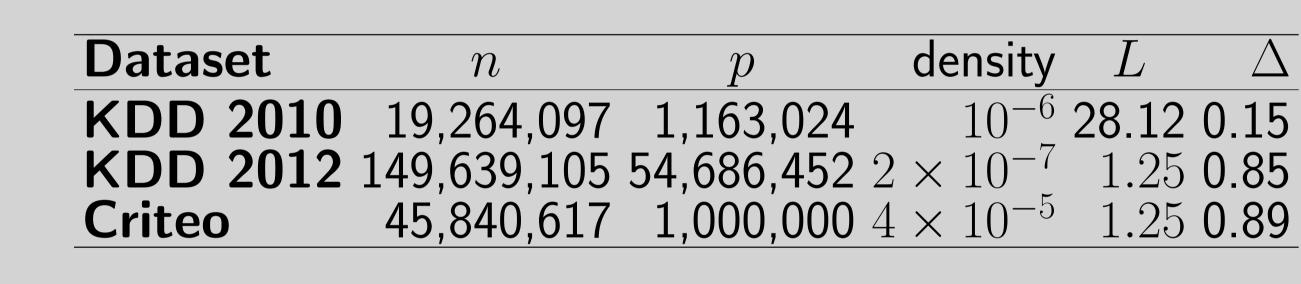
In both cases, the convergence rate is the same as Sparse Proximal SAGA \Longrightarrow ProxASAGA is **linearly faster** up to constant factor. In both cases the **step size does not depend on** τ .

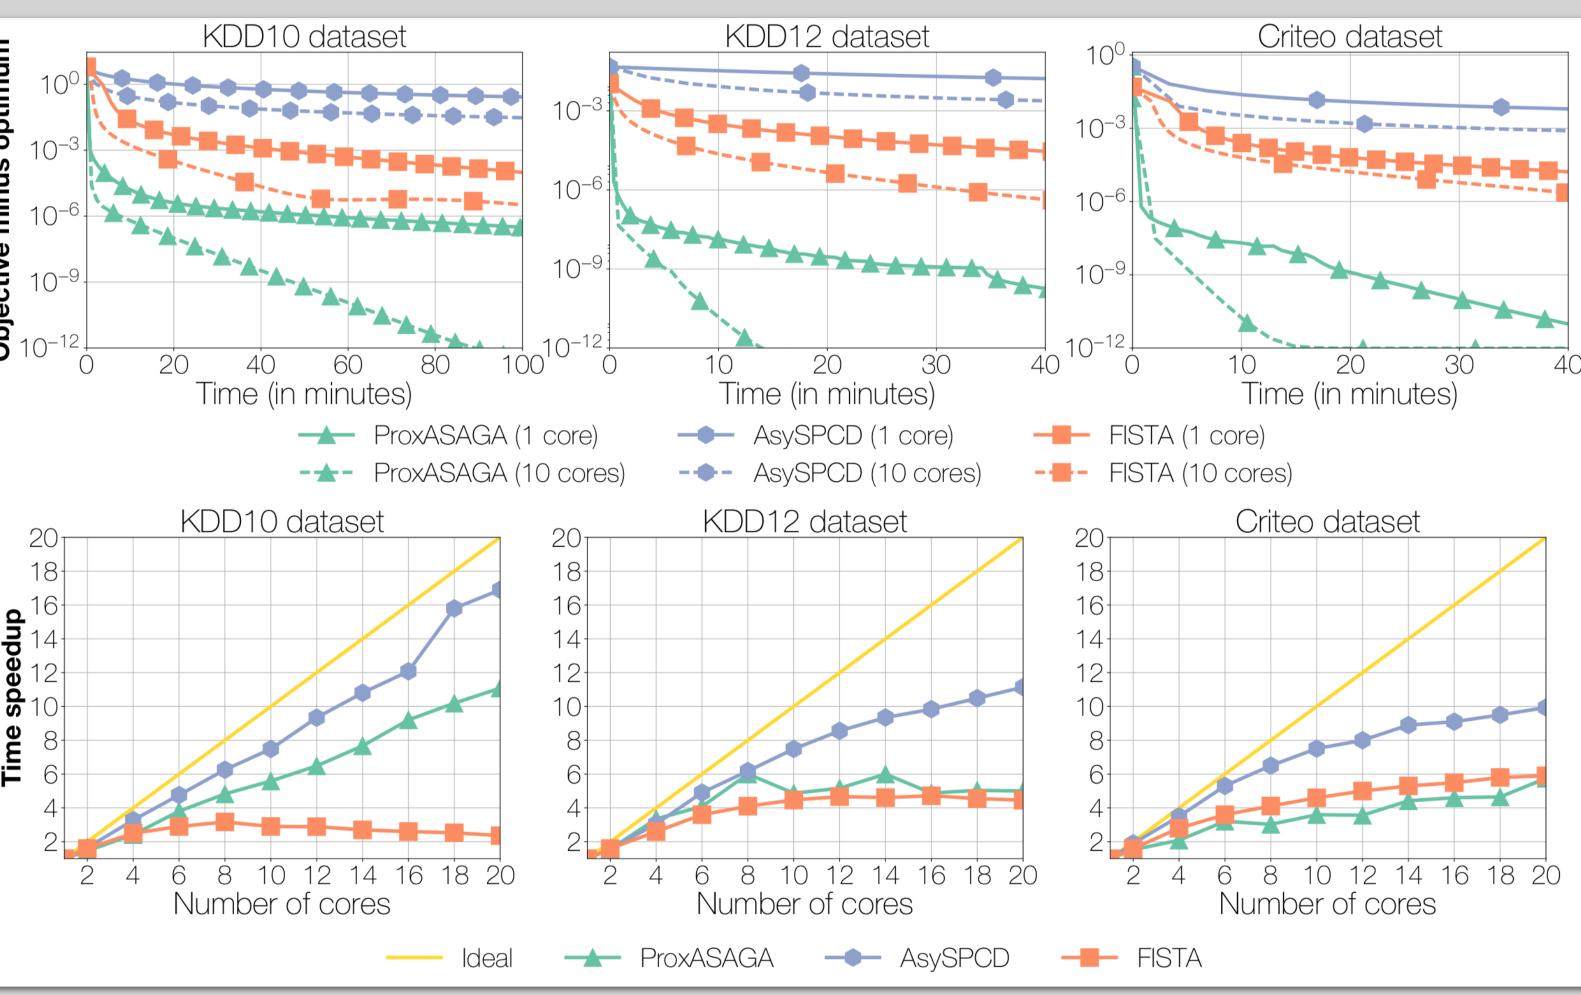
If $\tau \leq 6\kappa$, a universal step size of $\Theta(1/L)$ achieves a similar rate than Sparse Proximal SAGA, making it adaptive to local strong convexity (knowledge of κ not required).

Experimental results

Comparison on 3 large-scale datasets on an elastic-net regularized logistic regression model:

$$\min_{x} \min_{t} \sum_{i=1}^{n} \log \left(1 + \exp(-b_i \boldsymbol{a}_i^{\mathsf{T}} \boldsymbol{x})\right) + \frac{\lambda_1}{2} \|\boldsymbol{x}\|_2^2 + \lambda_2 \|\boldsymbol{x}\|_1 \quad ,$$





Highlights: ProxASAGA significantly outperforms existing methods, significant speedup (6x to 12x) over the sequential version.

References

- 1. Niu, F., Recht, B., Re, C. & Wright, S. Hogwild: A lock-free approach to parallelizing stochastic gradient descent. in NIPS (2011).
- 2. Mania, H. *et al.* Perturbed iterate analysis for asynchronous stochastic optimization. *SIAM Journal on Optimization* (2017).
- 3. Leblond, R., Pedregosa, F. & Lacoste-Julien, S. ASAGA: asynchronous parallel SAGA. *AISTATS* (2017).
- 4. Defazio, A. et al. SAGA: A fast incremental gradient method with support for non-strongly convex composite objectives. in NIPS (2014).
- 5. Schmidt, M., Le Roux, N. & Bach, F. Minimizing finite sums with the stochastic average gradient. *Mathematical Programming* (2016).